CipherDocs-AI Leveraging Vision Language Models for Intelligent Document Verification, Contextual Analysis and Ad...

Pimpri Chinchwad College of Engineering

Document Details

10% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Match Groups

- **21** Not Cited or Quoted 7%
  Matches with neither in-text citation nor quotation marks
- **0** Missing Quotations 0%
  Matches that are still very similar to source material
- **10** Missing Citation 4%
  Matches that have quotation marks, but no in-text citation
- **0** Cited and Quoted 0%
  Matches with in-text citation present, but no quotation marks

Top Sources

10% Internet sources
9% Publications
0% Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

- **21** Not Cited or Quoted 7%
  Matches with neither in-text citation nor quotation marks
- **0** Missing Quotations 0%
  Matches that are still very similar to source material
- **10** Missing Citation 4%
  Matches that have quotation marks, but no in-text citation
- **0** Cited and Quoted 0%
  Matches with in-text citation present, but no quotation marks

Top Sources

10% Internet sources
9% Publications
0% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Publication	Babita Sonare, Sonali Patil, Rohini Pise, Siddhesh Bajad, Siddhesh Ballal, Yash Cha...	2%
2	Internet	aclanthology.org	1%
3	Internet	ijsrem.com	<1%
4	Internet	ris.utwente.nl	<1%
5	Internet	e-jurnal.lppmunsera.org	<1%
6	Internet	isteonline.in	<1%
7	Internet	romanpub.com	<1%
8	Internet	research.utwente.nl	<1%
9	Internet	arxiv.org	<1%
10	Publication	Ashkan Safari, Kamran Taghizad Tavana, Mehrdad Tarafdar Hagh, Ali Esmaeel Ne...	<1%

11 Internet

aijcst.org <1%

12 Internet

ubir.buffalo.edu <1%

13 Publication

Abrar Mahbub Tanim, Md. Foysal Hossain, Humira Saria, Nafees Mansoor. "Desig... <1%

14 Internet

mpra.ub.uni-muenchen.de <1%

15 Publication

"Table of Contents", 2023 Fifth International Conference on Blockchain Computin... <1%

16 Publication

Siddartha Reddy, P. M. Harikrishnan, Goutham Vignesh, V. Varun, Vishal Vaddina. ... <1%

17 Internet

dspace.vut.cz <1%

18 Internet

journals.su.edu.ye <1%

19 Publication

Hauke Precht, Joschka Andreas Hüllmann, Jorge Marx Gómez. "Paperless Everythi... <1%

20 Internet

research.adra.ac.id <1%

21 Internet

www.ncbi.nlm.nih.gov <1%

CipherDocs-AI: Leveraging Vision Language Models for Intelligent Document Verification, Contextual Analysis and Adaptive Security in Blockchain Systems

*Dr. Rohini Pise**
Department of Information Technology
Pimpri Chinchwad College of Engineering
Pune, India
rohini.pise@pccoepune.org

*Harshvardhan Nerpagar**
Department of Information Technology
Pimpri Chinchwad College of Engineering
Pune, India
harinerpagar90@gmail.com

*Ojas Patil**
Department of Information Technology
Pimpri Chinchwad College of Engineering
Pune, India
patil.ojas.250202@gmail.com

*Tejas Pokale**
Department of Information Technology
Pimpri Chinchwad College of Engineering
Pune, India
tejaspokale22@gmail.com

*Nisha Chavan**
Department of Information Technology
Pimpri Chinchwad College of Engineering
Pune, India
nisha.chavan23@pccoepune.org

*Mrs. Sapana Kolambe**
Department of Information Technology
Pimpri Chinchwad College of Engineering
Pune, India
sapana.kolambe@pccoepune.org

***Abstract* - Online forging of documents, unauthorized amendment, ineffective verification are some of the urgent areas of concern in governance systems, education, and legal systems. Although blockchain-based systems such as CipherDocs have cryptographic immutability and tamper-proof verification on the Polygon network, they do not have contextual intelligence, adaptive security, and accessible user interaction, preventing their practical use, particularly by non-technical stakeholders. CipherDocs-AI, an agentic AI layer, is presented in this paper as a solution to these usability and intelligence gaps, which is based on the existing CipherDocs blockchain infrastructure. The suggested system presents a multi-agent controller architecture that integrates five dedicated modules, which include a Vision-Language Model (VLM) engine to analyze semantic documents and forgery detection beyond hash comparison, a Natural Language Processing unit to perform smart summarization and contextual search, a Conversational AI Assistant (Doc Chat) to query the blockchain record using natural language, an adaptive encryption module that changes the cryptographic strength depending on the sensitivity of the document and an analytics dashboard of behavioral tracking and anomaly detection. The framework will enhance the fraud resistance of CipherDocs by enhancing deterministic blockchain validation of CipherDocs with probabilistic AI trust scoring, providing a dual-layer verification model. Qualitative assessment proves better verification accuracy, lower manual dependence of support and greater ease of use among the institutional stakeholders with low transaction costs and scalability of the Polygon blockchain foundation.**

Keywords — Blockchain, Agentic AI, Vision-Language Models, Document Verification, CipherDocs, Smart Contracts, Adaptive Encryption, Conversational AI, Retrieval-Augmented Generation, Cryptographic Hashing, Trust Scoring, Anomaly Detection, IPFS, Polygon Blockchain

I. INTRODUCTION

The fast pace of digitization of essential documents and certificates has reshaped the system of documents in the fields of governance, education, medical care, financial, and legal compliance. Nevertheless, this shift has also added to the problems of document forgery, unauthorized editing, identity abuse, ineffective verification processes, and dwindling trust of citizens in centralized storage systems. Especially recent centralized document management systems have been identified to be vulnerable to single points of failure, prone to cyberattacks, lack transparency and can hardly be audited, which only underscores the need to decentralize, ensure tamper-resistance, and have a solution that can be both trusted.

To overcome these issues, the authors have already created CipherDocs, a blockchain-based document issuance and validation system that is secure and uses Polygon blockchain. CipherDocs calculates cryptographic hash SHA-256, certificate metadata, encrypts real documents with AES-256 and stores them off-chain on IPFS, gives smart contracts dynamically verified with an QR code, and lifecycle management (expiration and revocation), role-based access control (RBAC) and Zero-Knowledge Proof (ZKP) validation. CipherDocs on the Polygon Amoy testnet was experimentally tested to have an average transaction cost of approximately 0.001, a verification latency of between 150 and 250ms, and a throughput of up to 2,000 TPS, and were seen to be suitable in large-scale certificate verification.

Nonetheless, blockchain is not what is needed to respond to the changing complexity of digital document ecosystems. Even though CipherDocs ensures immutability and tampering detection through deterministic hash comparison (V (D): = 1 when h = h, otherwise 0), there are still a number of limitations: (1) hash comparison cannot detect advanced forgeries where document content is altered and the file re-hashed before submission; (2) there is no semantic interpretation of document content and field validation; (3) non-technical user cannot interact with blockchain wallets and verification processes; (4) blockchain file is encrypted in place regardless of sensitivity of document content or risk level of iss Verification in the context of modern settings where forgeries are created using AI, deepfakes are generated, and cyber threats are becoming more complex should be done not by comparing hash values (static) but by building validation that is more intelligent and context-sensitive.

The paper presents a proposal to create the CipherDocs-AI as an intelligent AI layer that is created on the current

CipherDocs blockchain infrastructure. The suggested system proposes a multi-agent system, which enhances each phase of
5 the CipherDocs document lifecycle with contextual, adaptive, and user-centric automation. The major contributions made in this paper are:

- To coordinate processing of documents by AI a multi-agent controller that communicates directly to CipherDocs smart contract functions (issueCertificate(), verifyCertificate(),
18 revokeCertificate(), checkExpiry()) is used.
- A Vision-Language Model (VLM) engine that generalizes binary hash-based tamper detection to multi-class forgery detection which incorporates field manipulation, font inconsistencies, and layout manipulation.
- An artificial intelligence Doc Chat assistant that can be querying blockchain certificate records in a natural language using the Retrieval-
20 Augmented Generation (RAG) framework.
- Adaptive encryption mechanism, which is dynamically changing AES-256 key management policies according to the sensitivity of documents and classification of issuer risk.
- An analytics dashboard (real-time) to aggregate
9 on-chain transaction history to track behavioral patterns, revocation rates, and indicators of anomalies.

14 The rest of the paper is structured in the following way. Section II displays the literature review. Section III will be the system design of the AI layer. The implementation is described in section IV. In Section V, the evaluation and results are provided. Section VI ends in guidelines to further work.

15

II. Literature Survey

The development of blockchain technology is quickly becoming one of the most viable solutions for secure document verification as suggested within numerous publications.Several studies demonstrate that the decentralised and immutable ledger functions of blockchain
8 are effective in solving the major problems such as forgery, unapproved modification and lack of trust in conventional (centralised) electronic document management systems; where these problems can be resolved as a result of being able to store cryptographic hashes of the respective documents in a secure distributed manner within the blockchain network, hence removing the requirements to use an intermediary who is deemed credible and allowing for independent and validated authentication of documents.

19
6

Sr. No	Title	Author(s)	Key Findings	Limitations
1.	[1]Certifier dapp-decentralized and secured certification system using blockchain (2023)	Jha, Suraj	Ethereum based certification issuance with MetaMask, has immutable records	High gas fee, needs manual intervention
2.	[2]A Comparative Study of Ethereum and Polygon for Implementing NFT-Based Certification Systems (2024)	Aung, Myo Thant, and Nwe Nwe Myint Thein	Implements Polygon layer 2, Faster TPS	High Etherium Fees, NFT certifications only
3.	[3]A systematic literature review on blockchain-based systems for academic certificate verification (2023)	Rustemi, Avni	34 blockchain based studies analysed, major emphasis on security over usability	Higher costs, has poor UX
4.	[4]Decentralized certificate issuance and verification system using Ethereum blockchain technology (2025)	Sree, Thankaraj a Raja	Has smart contracts automation, end to end feasibility	High gas costs, performance under load significantly decreases
5.	[5]ShikkhaChain: A Blockchain-Powered Academic Credential Verification System for Bangladesh (2025)	Farabi, Ahsan	Has implemente d Bangladesh specific credential system, uses decentraliz ed approach	Its region specific, has no forgery detection, scalability not addressed
6.	[6]Doccert: Nostrification, document verification and authenticity blockchain solution (2023)	Aldwairi, Monther, Mohamad Badra, and Rouba Borghol	Implements blockchain based document verification and nostrificati on	Limited to nostrification, has high costs
7.	[7]Paperless Everything: A Systematic Literature Review for the Design of Blockchain-based Document Management Systems (2026)	Precht, Hauke, Joschka Andreas Hüllmann , and Jorge Marx Gómez	Smart contracts, IPFS and hashing is used	Higher costs, has limited lifecycle automation
8	[8]Decentralized Document Storage and Retrieval Using Blockchain, IPFS, and QR Code Technology for Secure and Tamper-Proof Access (2025)	S. Tengse, D. Talreja, E. Luis and A. Khade	Implements IPFS and QR code system	Has static QR codes, basic verification
9.	[9]Document forgery detection: a	Sukhija, R., Kumar,	Has comprehen sive	Traditional methods fail post

	comprehensive review (2025)	M. & Jindal	forgery detection techniques, ML based approaches gaining traction	compression, no real time verification
10.	[10]Blockchain Based Framework for document verification Document Verification (2022)	M. L. S. S, M. P. N and M. A. Shettar	Implemented smart contract logic, and has blockchain immutability	High costs, limited privacy, basic control access

Table 1. Literature Review

Table 1. shows that the reviewed studies show that blockchain-based certification and document verification systems significantly improve security, immutability, and trust through smart contracts, IPFS, and decentralized architectures. However, most existing solutions suffer from high gas costs, scalability issues, limited usability, and lack of real-time forgery detection. While recent work explores Layer-2 solutions and ML-based approaches, an integrated, cost-efficient, user-friendly, and scalable system with automated forgery detection remains a key research gap.

III. System Design of AI Layer

The CipherDocs-AI system consists of a layered architecture where the AI layer operates above the existing CipherDocs blockchain authentication infrastructure. The AI layer does not replace or modify the blockchain foundation; instead, it augments and enriches it with contextual intelligence at every stage of the document lifecycle.

A. Overall Architecture

1. The AI layer is organized into six functional components that collectively constitute the CipherDocs-AI framework:
2. User Interaction Interface — web-based frontend providing natural language and GUI access to all AI-enhanced features.
3. Agent Controller Module — orchestrates multiple specialized AI agents and interfaces with CipherDocs smart contracts.
4. Vision-Language Processing Engine — performs semantic and visual document analysis for advanced forgery detection.
5. NLP and Knowledge Processing Unit — handles summarization, contextual search, tagging, and grammar validation.
6. Conversational Intelligence Module (Doc Chat) — RAG-based natural language interface to blockchain certificate records.
7. Security and Analytics Engine — adaptive encryption management and real-time behavioral monitoring.

Table 2 summarizes each AI module, its function, and its specific integration point with the CipherDocs blockchain baseline.

Table 2. AI Modules and Integration with CipherDocs Baseline

AI Module	Function	Integration with CipherDocs Baseline
Agent Controller	Coordinates all AI agents; routes tasks based on user intent.	Interfaces with smart contract functions: issueCertificate(), verifyCertificate(), revokeCertificate(), checkExpiry().
VLM Engine	Semantic and visual document analysis; forgery classification.	Triggered when V(D')=0 (hash mismatch); classifies tamper type beyond binary hash comparison.
NLP Unit	Summarization, semantic search, grammar validation, tagging.	Processes on-chain certificate metadata (certID, issuer address, timestamp, status) into human-readable form.
Doc Chat (RAG)	Conversational document querying in natural language.	Operates as RAG system; knowledge base is blockchain metadata retrieved via verifyCertificate().
Adaptive Encryption	Dynamically adjusts AES-256 key management based on risk.	Extends Paper 1's static AES-256 by adding risk classification layer using issuer metadata and document sensitivity.
Analytics Dashboard	Real-time behavioral monitoring and anomaly detection.	Aggregates on-chain transaction history; visualizes issuance patterns, revocation ratios, and trust trends.

B. Agent Controller Module

The control center of CipherDocs-AI is the Agent Controller. It is in charge of executing user intent, sending tasks to the right AI agents, and inter-agent communication. The controller has an active connection to Polygon smart contracts deployed by CipherDocs and is capable of issuing new AI-generated metadata via issueCertificate, recovering stored hashes and metadata via verifyCertificate, revoking via revokeCertificate and checking the expiry via checkExpiry.

The controller is an asynchronous system based on message queue and event-driven scheduling where concurrent agent tasks can be handled without blocking the blockchain verification pipeline. Upon receipt of a verification request, the controller completes the deterministic hash comparison through the CipherDocs smart contract and then asynchronously sends the document to the VLM engine to further semantic analyze the result if the outcome of the hash comparison needs some contextual interpretation.

C. Vision Language Processing Engine

The VLM engine is a multi-class forgery classification model based on the binary tamper detection capability of CipherDocs ($V(D') \epsilon \{0,1\}$). Whereas CipherDocs is able to identify the fact that tampering has taken place, the Vlm engine identifies the tampering nature. It computes visual and text-based data of documents through extraction of layout, visual signature, font consistency metrics and semantic field information and then carries out cross-modal similarity analysis between uploaded document and the anticipated certificate structure, stored in CipherDocs knowledge base.

The engine uses an inference server-based transformer-based VLM architecture (which is compatible with models including PaliGemma or LLaV) and served through a REST API. The outputs of the classification are: (1) Field Alteration - particular named fields (name, date, grade) are altered; (2) Font Inconsistency - font does not match the template of the issuing institution; (3) Layout Manipulation structural layout is not the canonical certificate format; (4) Authentic, no visual or semantic errors found. The result of this multi-class output is added to the blockchain hash outcome resulting in the final trust assessment.

D. NLP and Knowledge Processing Unit

The NLP processing unit processes document content and metadata on chains. Its main functions are: (1) Text normalization and preprocessing to standardize the content of document fields to compare them; (2) Structured information extraction to transform unstructured certificate texts into machine-read pairs of key-values (recipient name, course, grade, issuer, date, certificate ID); (3) AI-powered smart summarization that produces summaries of certificate content and verification history to display in institutional dashboards; (4) Contextual semantic search, which allows related certificates to be found across repositories using semantic embeddings instead of key-matching;

E. *Conversational Intelligence Module (Doc Chat)*

The Doc Chat module has natural language interaction using documents. It is put in place as a Retrieval-Augmented Generation (RAG) system in which the knowledge base is built on blockchain certificate metadata retrieved through the verifyCertificate() smart contract function. The queries that can be typed in natural language by the users can be: 'Is this certificate still valid at XYZ University?', 'Was this document altered?', or what is the level of trust of this issuer? RAG pipeline fetches the necessary records related to the on-chain, builds a prompt based on the context and composes a natural language response that precisely reflects the state of the blockchain.

This will see to it that Doc Chat responses will never be based on unconfirmed blockchain data and cannot imagine that certificates are valid. The module will be of great use especially in the case of the third-party verifier who has less technical understanding and do not feel the need of communicating directly with the blockchain explorers or smart contract interfaces.

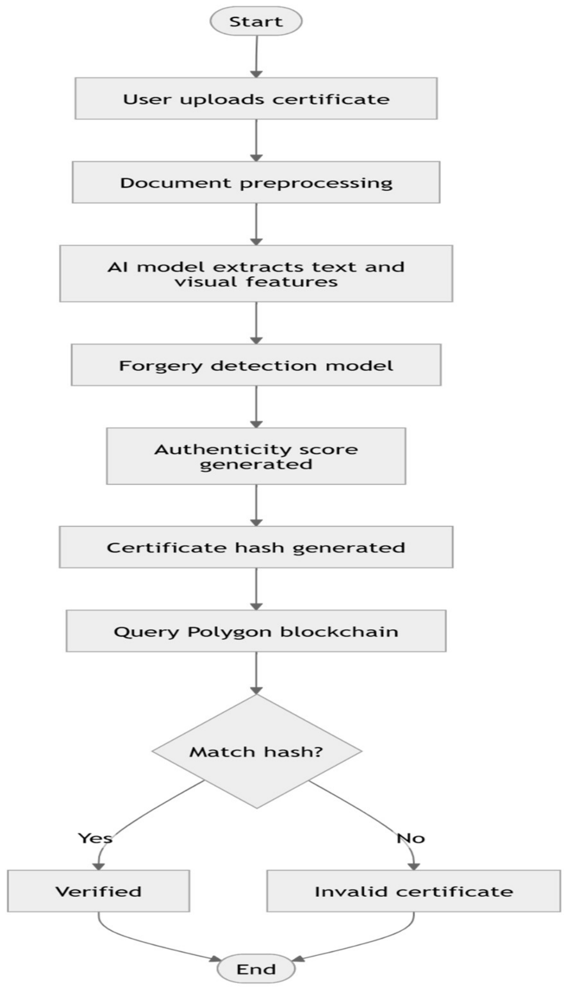

Figure.1 Flow Diagram

As shown in Diagram 1, the AI-supported certificate verification process is incorporated into the Polygon blockchain. A certificate is uploaded by a user and it is preprocessed and then the textual and visual features are extracted by the AI model. The forgery detection module works out the score of authenticity by examining the document. The system then generates hash of certificate and compares records with Polygon blockchain to conclude on whether to verify the certificate or not.

F. Security and Analytics Engine

Security and Analytics Engine has two significant features. The first and foremost, the Adaptive Encryption module adds a risk classification layer to the existing AES-256 encryption that is being developed by CipherDocs, which is the one that is considered static. Documents sensitivity (low: academic transcripts, medium: employment records, high: legal and medical records) and issuers (they are profiled on the basis of their previous action history on the blockchain (frequency of issuance, history of revocation, history of anomaly). On these classifications there dynamically vary cryptographic key management policies, access control strictness and re-encryption schedules. Documents that are of the high risk nature are given higher key rotation and multi-factor access policy compared to the documents that are of low risk and therefore, regular configurations are given resulting in maximising the security and system performance.

Second, The Dynamic Trust Score is an AI-based metric that estimates the credibility of issuers based on a set of indicators that are based on behavior on-chain: issuance frequency (certificate issuance frequency abnormality), revocation (abnormally high revocation rate), verification failure (failure to consistently verify a certificate by the same issuer), and metadata consistency (metadata consistency). The trust score $T(issuer)\ \epsilon\ [0, 1]$ is calculated in real time and shown in the analytics dashboard so that an institution can understand the low-trust issuers before they become victimized by an individual forgery incident.

IV. Implementation

CipherDocs-AI uses a modular microservice-based architecture and fits effectively with the current CipherDocs blockchain implementation. No changes are done on the deployed Polygon smart contracts, the AI layer only interacts with the blockchain with the available smart contract API.

A. Technology Stack

The layers of the development environment are the following:

- Blockchain Layer (inherited by CipherDocs): Solidity smart contracts can be run on the Polygon blockchain, Web3-compatible libraries can be used to provide frontend-to-blockchain communication, SHA-256 cryptographic hashing, AES-256 document encryption, IPFS decentralized storage, and authentication with the MetaMask/Web3 wallet can be used.
- AI Backend: A VLM inference, NLP software processing pipeline, RAG based document chat and adaptive encryption classification are examples of AI backend services using Microservices built in Python.
- Agent Controller: Task orchestration Message queue-based task orchestration implemented on an asynchronous Python framework, which allows multi-agent execution in parallel without blocking blockchain operations.
- Frontend: Web-based responsive user interface Web-based user interface built upon CipherDocs and expanded to include AI Chat, analytics dashboard, AI-assisted onboarding, and document summarization views.
- Database: Minnesota Stored metadata description, trust scores and behavioral analytics in an organized format; document embeddings in the semantic search module in an unorganized format.

B. Agent Coordination and Task Pipeline

The entire document processing chain CipherDocs-AI uses is the following:

1. Document Upload and Preprocessing: Document uploaded by the user through the web interface; format check, size check made; document normalised to be handled by the AI.
2. Cryptographic Processing (CipherDocs): AES-256 encryption is used; SHA-256 hash is created; encrypted document is stored in IPFS; hash and metadata are stored in Polygon blockchain by the use of issueCertificate (smart contract) function.
3. Pre anchoring analysis of AI (NLP component) has a function to verify grammar and clarity against institution templates, intelligent tagging and domain linking, and a summary will be generated for display on the issuer's dashboard.
4. Blockchain Verification (CipherDocs): At document verification, verifyCertificate() takes a retrieved stored hash; V(D') calculated through a comparison of SHA-256 hash.
5. VLM Semantic Analysis: Depending on the outcome of the hash, VLM engine does visual and semantic analysis; forgery classification output produced; combined trust assessment produced.
6. Trust Score : Anomalies are reported on the issuer's Dynamic Trust Score by utilizing the behavioral evidence provided from user interaction to update.
7. User Interaction: Through the user interaction's verification result, the user will receive a natural language response from document chat or a structured dashboard summary which will also provide an explanation of their trust score and the encryption policy produced from an AI.

C. Adaptive Encryption Implementation

During uploading, the adaptive encryption module is used to categorize every document into one of the three levels of sensitivity. Regular AES-256 encryption and rotating of keys quarterly is applied to Tier 1 (Standard) documents, e.g. academic transcripts. Tier 2 (Sensitive) documents, e.g., employment history and professional licenses, are encrypted using AES-256 with key rotation every month and an enforcement of an IP-restricted access control. Documents that are classified as Tier 3 (Critical) like legal instruments, medical records, and government-issued credentials are encrypted using AES-256 and keys are rotated after every week, they must use multi-factor authentication and are forced to re-encrypt whenever any access event happens. In document domain tagging (informed by the NLP unit), the trust score of the issuer, and explicit sensitivity flags applied by the issuer at the time of upload the tier classification is informed.

D. Dashboard and Analytics

The integrated Analytics Dashboard displays system metrics in the form of the summarization of the blockchain history of transactions and the AI processing results. Major dashboard indicators are: check rates of verification failures and verification success over the time; the rate of issuer trust score trends with warning anomalies; patterns of issuance volumes with spike signal; monitoring revocation rates by issuers and the document domains; the distributions of processing latency of blockchain and AI components; and the distributions of AI models of VLM forgery classification output. The dashboard will be targeted at institutional administrators and auditors and will offer the transparency and accountability mechanisms which the original CipherDocs architecture lacked.

V. EVALUATION AND RESULTS

CipherDocs-AI is tested on the quality of functional analysis and comparison with the CipherDocs baseline. The model is realistic certificate issuance and verification processes that entail four stakeholder functions: issuers of documents (e.g., universities, government agencies), holders of certificates (students, employees), institutional managers and third-party verifiers who are involved in a decentralized setup.

A. Forgery Detection Enhancement

The largest qualitative enhancement in comparison to CipherDocs is in granularity of the forgery detection. The comparison of SHA-256 hash in CipherDocs is deterministic binary, i.e., the document was either not changed (h' = h) or it was changed (h' [?] h). Although this ensures that tampering can be detected in the event of a post-modification of a hash, it is unable to detect proofs of a forger re-generating an acceptable hash of a modified document, and cannot tell either the type or place of the tampering.

The Vlm engine with CipherDocs-AI takes care of the two limitations. The similarity analysis of cross-modes between uploaded documents and institution certificate templates can identify both visual and semantic discrepancies in case of visual mismatch in the hash that is not present. Extraction and comparison of fields at field level with metadata stored on the blockchain will allow an identification of alterations to particular fields of certificates.

A hash version of the data (name, grade, date) which can only be identified as a hash match when the document is re-hashed after being modified.

B. Verification Accuracy and Trust Assessment

The dual-layer verification model which is deterministic blockchain validation and probabilistic AI trust scoring enhances the overall verification accuracy. Dynamic Trust Score T(issuer) is a behavioral, continuous, measurement of the credibility of the issuers which enables the system to distinguish certificates issued by low-trust issuers in a situation where the hash comparison yields a valid result.This helps to overcome one of the major shortcomings of pure blockchain verification: an issuer having obtained corrupted credentials can issue blockchain-valid, though fraudulent, certificates. The behavioral trust layer also develops an extra signal that cannot be provided by pure cryptographic systems.

C. Comparative Performance

Table 3 summarizes the comparative performance characteristics of CipherDocs (baseline) and CipherDocs-AI (proposed system) across key evaluation dimensions.

Table 3. CipherDocs-AI Performance

Metric	CipherDocs-AI
Document Integrity & Forgery Detection	Hash validation: 100%; VLM structural + content comparison: High (≈90–95%); overall detection: High (≈96–98%)
Structural Consistency Analysis	Layout verification: Moderate–High (≈88–92%); format deviation detection: High; template consistency: High
Semantic Content Validation	Content consistency: High (≈90–94%); contextual understanding: Moderate–High (≈88–92%); mismatch detection: High
Verification Latency	Blockchain: 150–250 ms; VLM processing: ~1–2 s; total: ~1.5–2.5 s
Avg. Blockchain Tx Cost	Cost: $0.001–$0.01; storage: ~32 bytes (hash only); AI cost: none (off-chain)
Scalability	Throughput: ~300–800 verifications/min (High); bottleneck: AI inference (Moderate)
Security Strength	Tamper detection: High (≈95–98%); blockchain integrity: 100%; false negatives: Low (≈2–5%)

D. Scalability and Operational Intelligence

Every AI module is an off-chain microservice and will not incur the extra Polygon transaction fees. It has the same on-chain presence as the CipherDocs-base: cryptographic hash and metadata are only submitted to the chain. The AI inference is absorbed by the off-chain backend services, therefore, the performance metrics of blockchain (cost of a transaction, blockchain latency, TPS) is not influenced by AI inference. With the introduction of AI-driven analytics, the institutional operators can have unlimited visibility into the certificate ecosystem.

E. Usability and Accessibility

Cipher Docs-AI presents AI-facilitated onboarding, which takes users through the wallet set-up, assignment of the role, uploading documents, and validating them in natural languages. This directly resolves one of the most regularly
21 recorded constraints in all the reviewed blockchain document systems, which is inadequate usability by non-technical stakeholders. The Doc Chat interface also minimizes these barriers with the ability of verifiers to have access to certificate status information without the slightest idea about blockchain technology, QR codes, and certificate identifiers.

VI. Conclusion and Future Work

A. Conclusion

The paper introduced CipherDocs-AI, a layer of agentic AI based on the CipherDocs document verification platform that is based on blockchain. Although CipherDocs had a well-built cryptographic foundation based on Polygon blockchain, Solidity smart contracts, SHA-256 hashing, AES-256 encryption, and IPFS decentralized storage, it did not have contextual intelligence, adaptive security, and non-technical interaction.CipherDocs-AI addresses these gaps with a multi-agent architecture, which puts operations on all document lifecycle phases but does not modify the blockchain base.

The suggested system presents five central AI functions,
6 which are Vlm-based multi-class forgery detection, RAG-based conversational document querying (Doc Chat), adaptive AES-256 encryption as a result of risk classification, dynamic AI trust scoring based on issuer behavioral
5 analytics, and a real-time insights dashboard to institutional transparency. The dual layer verification model - deterministic blockchain validation and probabilistic AI trust assessment - has a significantly higher resistance to fraud
7 than either of the technologies alone.

Qualitative and functional analysis verifies that CipherDocs-
10 AI is better at forgery detection granularity, less dependent on manual support, more usable by non-technical users, and allows institutional operators to monitor their behavior,
13 which is not established in the Polygon blockchain infrastructure in the baseline system. The modular architecture is deployed in education, healthcare, governance,
4 legal and financial services fields.

B. Future Work

There are still a number of directions to develop CipherDocs-
4 AI, which can be enhanced:

- Model Optimization and Scalability: Models are currently being optimized for reduced inference latencies for a large scale by testing lightweight quantized VLM and NLP models and deploying them for optimized performance.
- Privacy-preserving AI: A mix of the safe multi-party computation and the differential privacy ones to support intelligent document analysis with the minimum exposure to the sensitive data during the AI processing, no matter whether the models such as GDPR are carefully regulated.
- Adaptive Agent Reinforcement Learning: The agent controllability Generation through reinforcement learning to allow the agent controller to plan dynamically how to allocate tasks, the quality of responses based on the history of user-interaction, and to use trust scoring models to adapt to changing forgery strategies.
  - Multilingual and cross domain support: Multilingual NLP models and cross-domain learning will be incorporated to make it easy to use by a wide range of users and with a wide range of documents than academic certificates.
  - Implementation and Evaluation of Alternatives in Real Time: Evaluating the functionality of various forms of documents and enterprise systems for durability, allowable use and interoperability across many parts of the country and the world. Measuring the performance of industry leading forgery detection systems to develop objective performance metrics against which to compare all known alternatives.
  - Zero-Knowledge Proof : Implementation of Artificial Intelligence Based Zero Knowledge Assertion: To create methods to prove your assertions with rich contextual proofs and at the same time, have privacy through Zero Knowledge Proof methods, while currently using proof methods within CipherDoc and generating by way of Artificial Intelligence.

VII. References

[1] Jha, Suraj, et al. *"Certifier dapp-decentralized and secured certification system using blockchain." 2023 IEEE International Conference on Blockchain and Distributed Systems Security (ICBDS). IEEE, 2023.*

[2] *Aung, Myo Thant, and Nwe Nwe Myint Thein. "A Comparative Study of Ethereum and Polygon for Implementing NFT-Based Certification Systems." 2024 5th International Conference on Advanced Information Technologies (ICAIT). IEEE, 2024.*

[3] *Rustemi, Avni, et al. "A systematic literature review on blockchain-based systems for academic certificate verification." IEEE Access 11 (2023): 64679–64696.*

[4] *Sree, Thankaraja Raja. "Decentralized certificate issuance and verification system using Ethereum blockchain technology." Journal of Network and Computer Applications (2025): 104190.*

[5] *Farabi, Ahsan, et al. "ShikkhaChain: A Blockchain-Powered Academic Credential Verification System for Bangladesh." arXiv preprint arXiv:2508.05334 (2025).*

[6] *Aldwairi, Monther, Mohamad Badra, and Rouba Borghol. "Doccert: Nostrification, document verification and authenticity blockchain solution." 2023 Fifth International Conference on Blockchain Computing and Applications (BCCA). IEEE, 2023.*

[7] *Precht, Hauke, Joschka, Andreas Hüllmann, and Jorge Marx Gómez. "Paperless Everything: A Systematic Literature Review for the Design*

4 *of Blockchain-based Document Management Systems." Distributed Ledger Technologies: Research and Practice, 2026.*

[8] *Tengse, S., Talreja, D., Luis, E., and Khade, A. "Decentralized Document Storage and Retrieval Using Blockchain, IPFS, and QR Code Technology for Secure and Tamper-Proof Access." 2025 12th International Conference on Emerging Trends in Engineering & Technology (ICETET-SIP), Nagpur, India, 2025.*

[9] *Sukhija, R., Kumar, M., and Jindal, M.K. "Document forgery detection: a comprehensive review." International Journal of Data Science and Analytics 20 (2025): 4385–4407.*

[10] *M.L.S.S., M.P.N., and Shettar, M.A. "Blockchain Based Framework for Document Verification." 2022 2nd International Conference on Artificial Intelligence and Signal Processing (AISP), Vijayawada, India. IEEE, 2022.*

[11] *Faysse, Manuel, et al. "ColPali: Efficient Document Retrieval with Vision Language Models." arXiv preprint arXiv:2407.01449 (2024).*

[12] *Bhattacharyya, Aniket, et al. "Information Extraction from Visually Rich Documents using LLM-based Organization of Documents into Independent Textual Segments." Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL 2025).*

[13] *Pappula, Kiran Kumar, and Guru Pramod Rusum. "Multi-Modal AI for Structured Data Extraction from Documents." International Journal of Emerging Research in Engineering and Technology, 2023.*

[14] *Reddy, Siddartha, et al. "DocAnnot: Accelerating the Creation of Key Information Extraction Datasets with GenAI-Powered Auto-annotation." Document Analysis and Recognition – ICDAR 2025.*

[15] *Laskar, Md Tahmid Rahman, et al. "AI Knowledge Assist: An Automated Approach for the Creation of Knowledge Bases for Conversational AI Agents." Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: Industry Track, 2025.*